

The AI Workforce Blueprint

How to set up AI agents for every part of your business, step by step: projects, agents, skills, connectors and routines.

Operations · Support · Marketing · Sales · Finance · Development

Before you start

For more than 25 years I built and ran operations teams. As COO of a business coaching company, I helped grow revenue from about \$5 million to \$16 million with a team of more than 50 people. The thing that made that growth possible was never one brilliant hire. It was the system around the people: clear job descriptions, written procedures, the right access to the right tools, a weekly rhythm, and a manager who checked the work.

AI agents need exactly the same system. Most people skip it. They open a chat window, type a request, get a decent answer, and start over from zero the next day. That's not a workforce. That's a temp you re-train every morning.

This guide shows you how to set up AI the way you'd set up a team: one place for the work, a job description for every role, written procedures the AI follows every time, access to your tools, and a schedule for the recurring jobs. Then it walks you through what to hand off in each part of the business: operations, customer support, marketing, sales, finance, and development.

How to use this guide

- **Read Chapter 1 first.** It explains the six building blocks. Every other chapter uses them.
- **Then build one department, not six.** Pick the area that eats the most of your week and set that up first. Add the next one only when the first is running.
- **Do the "Your move" at the end of each chapter.** Each one takes 15 to 45 minutes.
- **Fill in every [bracket]** in a template before you use it.

Which tool this guide uses

The walkthrough uses **Claude** (from Anthropic), because as of October 2026 it has all six building blocks in one place: projects with agents that work in parallel, skills, connectors to your business apps, and scheduled routines. Screens and plan details change often, so where I name a menu or a plan, I've checked it against Anthropic's own documentation in October 2026. If something looks different when you get there, the idea still holds.

The method itself works in any serious AI platform. ChatGPT, Gemini and Copilot have their own versions of most of these pieces under different names. Learn the six building blocks and you can set up a workforce anywhere.

The one rule

Agents do the work. You own the outcome. An AI workforce is staff, not a vending machine. You decide what it's allowed to do on its own, you review what it produces, and you're accountable for anything that reaches a customer, a bank account, or a public page.

Chapter 1: The six building blocks of an AI workforce

If you've ever hired and managed people, you already know how this works. Every part of an AI workforce has a direct match in a human team.

Human team	AI workforce	What it does
A department	Project	One home for one stream of work, with its own files, history and settings
Job description	Instructions	Standing rules every agent in the project reads before it starts
Institutional knowledge	Memory	Facts and decisions the AI keeps so you never repeat yourself
Staff members	Agents (threads)	Workers that each take one task and report back when it's done
Standard operating procedures	Skills	Written step-by-step procedures the AI follows the same way every time
System access	Connectors	Secure links to your email, calendar, drive, CRM, books and more
The recurring calendar	Routines	Jobs that run on a schedule without you starting them

Here's how they fit together in Claude.

Projects: one department, one home

A project is one ongoing conversation where Claude acts as a **coordinator**. You drop work into it as it comes up, the way you'd send a message to a manager: "Answer these three support emails," "Draft next week's newsletter," "Reconcile September's expenses." The coordinator decides what each request needs and starts a **thread** for each piece of work. Threads run in parallel, each one does its task, and each one reports back. An **Overview** pane shows you which threads are done, which are still working, and which are waiting on you.

That's the big shift from a chat window. You stop doing the coordinating yourself.

Instructions and memory: train once

Project instructions are the brief every new thread starts from (up to 16,000 characters). Memory is the set of notes Claude keeps about the project as it works: decisions you made, preferences you stated, mistakes it shouldn't repeat. When you correct a thread, tell it "remember that," and every future thread starts with the correction.

Skills: your procedures, written down

A skill is a folder with a file called `SKILL.md` that teaches Claude how to do one job your way: how you write a refund reply, how you format an invoice summary, how you run a weekly report. Claude loads a skill automatically when the task matches it. Skills are where your know-how lives, so your workforce gets better every time you write one.

Connectors: access to your tools

Connectors let Claude read from and act in the apps you already use: Gmail, Google Calendar, Google Drive, Slack, QuickBooks, Stripe, HubSpot, Notion and many more. You connect each one once on your account, and your agents can use it.

Routines: the recurring calendar

A routine is a saved job (a prompt, the tools it can use, and a schedule) that runs on its own: every weekday morning, every Monday, the first of the month. Routines run in the cloud, so they keep working when your laptop is closed.

Your move

On a sheet of paper, draw your business as an org chart with six boxes: Operations, Customer Support, Marketing, Sales, Finance, Development (cross off any you don't have). Under each box, write who does that work today. If the answer is "me" in four or more boxes, this guide is for you.

Chapter 2: Map the work before you hire

You wouldn't hire someone without knowing what you need them to do. Same here. The fastest way to waste money on AI is to automate the wrong thing.

Step 1: List everything you do in a week

Prompt: Your weekly work inventory

I run a [type of business] with [number] people. Here is everything I can remember doing in the last two weeks, in no particular order: [brain-dump every task, big and small]. Organize this into a table with columns: Task, Department (Operations, Support, Marketing, Sales, Finance, Development, Admin), How often (daily, weekly, monthly), Minutes each time, and Minutes per month. Sort by minutes per month, highest first. Flag any task that appears more than once under different names.

Step 2: Score each task for handoff

Not every task is a good first job for an agent. The best first jobs are frequent, follow a pattern, use information that already exists somewhere, and are low-risk if the first draft is wrong.

Prompt: Score the handoff

Here is my task table: [paste]. For each task, score 1 to 5 on: Repeatable (does it follow the same steps each time?), Information ready (does the information live in email, docs or an app?), Low risk (could a wrong first draft be caught before it does harm?). Add the three scores. Then sort by total score, and for the top 10 tell me in one line what an AI agent would do and what I would still need to check.

Step 3: Pick your first three jobs

From the top of the list, pick three jobs that:

1. Each take at least an hour a month.
2. Live in the same department, so you set up one project, not three.
3. Have at least one where you can see results within a week.

Those three jobs are your first hires. Everything in the next four chapters, you'll set up for them.

What to keep for yourself (for now)

Keep these human until your workforce has earned trust: anything that sends money, anything that signs or changes a contract, anything that deletes data, anything that goes to a customer during a crisis, and hiring or firing decisions. Agents can prepare all of these. A person should press the button.

Your move

Run both prompts and write your three first jobs at the top of a page. Next to each, write the department and how many minutes a month it costs you today. That number is your baseline for measuring results.

Chapter 3: Set up your headquarters

Now you build the home for your first department. This takes about 30 minutes.

Step 1: Check your plan

As of October 2026, Claude offers Free, Pro, Max, Team and Enterprise plans. Projects with coordinating agents (the version with threads, described in Chapter 1) are in public beta on **Pro and Max** plans and are rolling out gradually; Anthropic says Team and Enterprise are coming. Skills work on every plan, and routines work on Pro, Max, Team and Enterprise. Check current prices at claude.com/pricing.

If you don't see Projects in your sidebar yet, you can join the waitlist from Anthropic's Projects documentation. Meanwhile, use a standard Claude chat project for each department, with the same instructions, skills and connectors from this guide. You'll do the coordinating yourself until the rollout reaches you, but everything you build carries over.

Step 2: Turn on the basics

1. In **Settings > Capabilities**, turn on **Code execution and file creation**. Skills need it.
2. Go to **Customize > Connectors** and connect the two or three apps your first three jobs need (more on this in Chapter 5).

Step 3: Create the project

1. Go to **claude.ai/code** (or the **Code** tab in the Claude desktop app) and select **Projects**, then **New project**.
2. **Name**: name it after the department, like "Customer Support" or "Marketing Engine."
3. **Goal (optional but worth it)**: one line describing the outcome, like "Answer every customer email within 4 business hours" or "Publish two posts and one newsletter every week."
4. **Context**: click **Add** and upload the files your agents should read: your price list, FAQ, brand guide, past newsletters, policies. You can also add a Google Drive folder. If your department works on a website or code, add the GitHub repository here too. Code isn't required: a project can work entirely on documents.
5. Click **Create project**.

Step 4: Write the job description

Open the gear icon (**Project settings**), then **Memory > Project instructions**. A good brief answers five questions: what is this project for, where does the work live, how does an agent check its own work, what should it do when something is missing, and what needs your approval first.

Template: Project instructions

This project runs [department] for [business name], a [one-line description of your business] serving [who your customers are].

The work: [list the 3 first jobs in one line each].

Voice: [3 to 5 rules: e.g., warm and direct, short sentences, no exclamation marks, sign off as "[name]"].

Sources of truth: prices and packages are in [file]; policies are in [file]; customer records are in [connector/app].

Before you call work done: re-read it against the relevant skill, check every number against the source, and list anything you assumed.

When something is missing: if you can't reach a file, an app or a fact you need, say exactly what's missing in your first message and stop. Don't guess or make up details.

Needs my approval first: sending anything to a customer, posting anything publicly, any refund or payment, deleting anything, and anything that changes a price.

Step 5: Choose models and pace

In **Project settings > General** you'll see **Thread model** and **Thread effort**. A new project runs every thread on Anthropic's most capable model at high effort, which uses your plan's allowance fastest. For routine writing and sorting work, a smaller model or lower effort is usually enough; keep the top setting for analysis and anything complicated.

Then tell the coordinator how you want it to work. It saves these preferences to memory:

- "Propose threads and wait for my go-ahead before starting them." (Use this for your first week.)
- "Run at most two threads at a time."
- "Only post when something finishes or is blocked."

Step 6: Send your first small job

Send one small, real piece of one of your three jobs. When the thread finishes, open it and read what it did. If it assumed something wrong, correct it and say "remember that." This one feedback loop is how your workforce gets trained.

Your move

Create your first project, paste in your instructions, and send one small task. Don't send a batch until you've seen one come back the way you want.

Chapter 4: Write your procedures as skills

Instructions say *who* the project is. Skills say *how* a specific job gets done. If you find yourself pasting the same directions into requests more than twice, that's a skill waiting to be written.

What a skill looks like

A skill is a folder containing a file named `SKILL.md`. The top of the file has a short header with a name and a description. The description matters most: it's how Claude decides when to use the skill. The rest is your procedure in plain language.

Template: SKILL.md

```
---
name: support-refund-reply
description: Use when replying to a customer who asks for a refund, a cancellation, or says a
purchase didn't work.
---
```

Goal: a reply the customer can act on in one read, in our voice, that follows our refund policy exactly.

Steps: 1. Find the order in [Stripe/your store] by the customer's email. Note the product, date and amount. 2. Check the policy: full refund within [30] days, no questions asked. After that, offer [alternative]. 3. Write the reply: thank them by first name, confirm what they bought, say exactly what happens next and when. 4. Keep it under 120 words. No apologies beyond one sentence.

Never: issue the refund yourself, promise a date you can't confirm, or blame the customer.

Example of a great reply: [paste one real reply you were proud of].

Let Claude write the first draft

You don't have to write skills from scratch. Describe the job and let Claude draft it.

Prompt: Turn a task into a skill

I want to create a Claude skill for this recurring task: [describe the task]. Here is how I do it today, step by step: [explain or paste notes]. Here are two examples of finished work I was happy with: [paste]. Write a complete SKILL.md file with a name (lowercase, hyphens), a one-sentence description that says exactly when to use it, the goal, numbered steps, a "Never" list, and one example. Ask me up to 3 questions first if anything is unclear.

Install the skill

You have two options:

1. **On your account (works everywhere, including chat):** put the `SKILL.md` in a folder with the same name as the skill, zip the folder, then go to **Customize > Skills**, click **+**, choose **+ Create skill**, then **Upload a skill**. Toggle it on.
2. **In a project with a repository:** save it at `.claude/skills/[skill-name]/SKILL.md` in that repository. Every thread in the project loads it automatically.

On Team plans you can share a skill with colleagues from the ... menu next to it in **Customize > Skills**. They turn it on themselves, and they get your updates automatically.

Skills worth writing first

- **Voice and style:** how you write, words you never use, how you open and close.
- **Your most common reply:** the email or message you send most often.
- **Your weekly report:** what goes in it, in what order, from which sources.
- **Your quality checklist:** what “done” means for your most important deliverable.

Your move

Pick the most repetitive of your three first jobs. Run the “Turn a task into a skill” prompt, read the result, fix anything that isn’t how you’d do it, and install it. Then give the project a task that should use it and check whether it did.

Chapter 5: Connect your systems

An agent that can't see your inbox, calendar or books has to wait for you to copy and paste. Connectors remove that bottleneck. They're also where most of the risk is, so set them up deliberately.

How to connect

1. Go to **Customize > Connectors** (claude.ai/customize/connectors).
2. Find the app and click **Connect**. You'll sign in to that app and approve access.
3. Every project thread can now use it. There's no per-project setup.

Which connectors each department needs

Department	Common connectors
Operations	Google Drive or Microsoft 365, Slack, Notion, Asana or Linear, Google Calendar
Customer support	Gmail or Outlook, your help desk, Stripe or your store, Slack
Marketing	Google Drive, Canva, your email platform, Notion, analytics
Sales	Gmail, Google Calendar, HubSpot or your CRM, Granola or your meeting notes tool
Finance	QuickBooks, Stripe, PayPal, Google Sheets, Gmail (for receipts)
Development	GitHub, Linear or Jira, Sentry, Slack

Availability changes often, so check the connector directory for your exact tools. If an app has no connector, you can usually work around it with exports: drop a weekly CSV into the project's files.

The three access rules

1. **Connect what the job needs, nothing more.** If your support agent only needs to read orders, it doesn't need your accounting system.
2. **Remove connectors from routines that don't use them.** When you create a routine, every connector is included by default, and a routine can use them without asking. Delete the ones it doesn't need (Chapter 6).
3. **Anything an agent does through a connector is done as you.** An email sent through your Gmail is from you. A record changed in your CRM shows your name. Treat connector access the way you'd

treat giving a new hire your password, which is why your instructions say what needs approval.

Your move

Connect only the apps your three first jobs need. Then ask your project: "List every connector you can use right now and one way each one helps with [department]." Remove anything you don't recognize or don't need.

Chapter 6: Put recurring work on autopilot with routines

Once a job works well when you ask for it, make it run on its own. A routine is a saved prompt plus the tools it may use, attached to a trigger: a schedule, a call from another app, or (for developers) an event in GitHub.

Create a routine

The easiest way is to ask from inside your project: "Turn the Monday report into a routine that runs every Monday at 7:10am." Claude sets it up and it appears on the project's **Routines** tab. You can also create one directly:

1. Go to **claude.ai/code/routines** and click **New routine**.
2. **Name** it and write the **prompt**. This is the most important part (see the template below).
3. Choose the **trigger**: hourly, daily, weekdays or weekly, or a one-off date. For something like "the 1st of every month," pick the closest option and ask Claude to set a custom schedule. Times are in your local time zone. Tip: pick a few minutes past the hour (7:10, not 7:00), since on-the-hour runs can start a few minutes late.
4. Under **Connectors**, remove every connector the routine doesn't need.
5. Click **Create**, then **Run now** once to test it.

The minimum interval for a schedule is one hour, and routines draw on your plan's usage the same way your own requests do.

Write a routine prompt that works unattended

A routine runs while you're not there, so the prompt has to stand on its own. Nobody can answer a follow-up question.

Template: Routine prompt

Job: [one sentence: what to produce].

Inputs: read [which connector, folder or file], covering [time window, e.g., the last 7 days].

Steps: [3 to 6 numbered steps; name the skill to use if there is one].

Output: [what to deliver and where: a file in the project, a draft email to me, a message in a Slack channel].

Done means: [how to know it's right, e.g., every number ties to the source; under 1 page].

Don't: send anything to customers, change any record, or delete anything. If an input is missing, report what's missing as the output and stop.

Check that it actually worked

A green status on a routine run means it started and finished without a technical error. It does *not* mean the job was done well. For the first few weeks, open each run and read what it did.

Your move

Choose the one job from your list that happens on the same day every week. Make it a routine, run it once with **Run now**, and read the result before you let it run on schedule.

Chapter 7: Department playbooks

Here's what to hand off in each part of the business, with the skills, connectors, routines and approval rules that make it work. Start with the department you chose in Chapter 2.

Operations

Hand off: meeting notes and action items, weekly status reports, SOP writing and updating, vendor research and comparisons, process mapping, inbox triage.

Skills to write: weekly status report format, SOP template, meeting-notes format.

Routine: *Monday 7:10am.* "Read last week's tasks in [project tool] and messages in [Slack channel]. Write a one-page status report using the status-report skill: done, in progress, blocked, decisions needed. Save it to the project and send me a summary in Slack."

Starter task:

Prompt: Turn a process into an SOP

Here's how we do [process] today, in my own words: [describe]. Write an SOP with: purpose, who owns it, trigger (when it starts), numbered steps, the tools used at each step, what "done" looks like, and the 3 most common mistakes and how to avoid them. Then list anything I left out that a new person would need to ask.

Needs approval: any change to a live process, anything sent to a vendor.

Customer support

Hand off: first drafts of replies, tagging and sorting the inbox, building and updating the FAQ, spotting trends in complaints, drafting help articles.

Skills to write: reply voice and rules, refund and cancellation reply, escalation rules (what goes straight to you).

Routine: *Weekdays 8:10am.* "Read new support emails since the last run. Sort them into: answer from FAQ, needs order lookup, needs me. For the first two groups, write draft replies using the support skills and save them as Gmail drafts. Send me a list of the 'needs me' emails with one line on each."

Starter task:

Prompt: Build the FAQ from your inbox

Read the last [90] days of customer emails in [inbox/label]. Group the questions into themes, count how often each comes up, and write an FAQ with the top 15 questions, each answered in under 80 words using our policies in [file]. Flag any question where our answer is unclear or our policy is inconsistent.

Needs approval: every reply that goes out, until the drafts are consistently right. Then let the routine send FAQ-only answers and keep approving the rest.

Marketing

Hand off: content calendars, repurposing one piece into many (a newsletter into posts, a call recording into an article), first drafts, SEO briefs, performance summaries.

Skills to write: brand voice, post formats for each channel, newsletter structure.

Routine: *Friday 9:10am.* "Using the content calendar in [file] and this week's notes in [folder], draft next week's [newsletter and 3 LinkedIn posts] using the brand-voice skill. Save drafts to the project's files and send me a summary."

Starter task:

Prompt: One idea, a week of content

Here is one idea or piece I've already written: [paste]. My audience is [who]. Turn it into: 1 newsletter (under 400 words), 3 LinkedIn posts with different hooks, 5 short social posts, and 1 email subject line test (3 options). Follow my brand-voice skill. Label each piece with the channel and the day of the week it should go out.

Needs approval: everything that's published. Always.

Sales and client acquisition

Hand off: prospect research, personalized first drafts of outreach, follow-up reminders, call preparation, proposal drafts, CRM updates after calls.

Skills to write: ideal client profile, outreach voice, proposal template, discovery-call prep format.

Routine: *Weekdays 7:40am.* "Check my calendar for today's sales calls. For each, research the company and person from [CRM and the web], and write a half-page prep sheet: what they do, likely problem, 3 questions to ask, and the offer that fits. Then list every deal in [CRM] with no activity in 7+ days and draft a follow-up for each."

Starter task:

Prompt: Call notes to proposal

Here are my notes from a discovery call with [name, company]: [paste]. Using my proposal skill and my packages in [file], draft a proposal: their situation in their words, the result they want, the recommended package, scope, timeline, price, and next step. Then write a short email to send with it.

Needs approval: every message to a prospect, every price.

(For a full client-getting system with AI, see [The AI Client Playbook](#) and [The AI Client Kit](#) on [virtualexpert.net](#).)

Finance

Hand off: categorizing transactions, matching receipts, monthly summaries, chasing unpaid invoices (drafts only), cash-flow snapshots, preparing questions for your accountant.

Skills to write: monthly finance summary format, expense categories and rules, collections email sequence.

Routine: *Monthly, on the 1st at 8:10am (a custom schedule).* "Pull last month's profit and loss and accounts receivable from [QuickBooks] and payments from [Stripe]. Write a one-page summary: revenue vs last month, top 5 expenses, anything unusual, invoices over 30 days late. Draft (don't send) a polite reminder for each late invoice. Save everything to the project."

Starter task:

Prompt: Find the money leaks

Here are my last [3] months of expenses: [connect your books or paste an export]. Group them by category and vendor, show the monthly trend, and flag: subscriptions I may not be using, costs that went up more than 10%, and duplicate charges. Rank the flags by how much money is at stake.

Needs approval: any payment, refund, transfer, or change to your books. Agents prepare; you or your bookkeeper post. AI summaries are not tax or legal advice, so check anything important with your accountant.

Development and website

Hand off: bug fixes, small features, website updates, code review, dependency updates, documentation, fixing failed checks.

Setup: add your GitHub repository to the project (it needs the Claude GitHub App installed). Each thread works on its own branch and opens a pull request, so nothing reaches your live site until you approve and merge it. Put rules about the codebase in a `CLAUDE.md` file in the repository; every thread reads it.

Skills to write: your test and check commands, your deployment checklist, your code review checklist.

Routine: a routine can also run when a pull request opens in GitHub: "Review this pull request against our code-review skill and leave comments." Or on a schedule: *Weekly:* "Check for outdated dependencies and open one pull request with safe updates."

Starter task: paste a bug report into the project. The thread finds the cause, fixes it on a branch, and opens a pull request for you to review.

Needs approval: merging to your main branch, anything that touches payments, customer data or production settings.

Admin and people (bonus)

Hand off: scheduling, travel research, onboarding checklists, job descriptions, interview question sets, policy drafts. Keep hiring decisions, performance conversations and anything personal and private in human hands.

Your move

Read the playbook for your first department. Write its skills, connect its apps, and set up one routine. Leave the other departments alone until this one runs for two weeks without surprises.

Chapter 8: Manage your AI workforce

A team without a manager drifts. So does an AI workforce. Here's the management rhythm that keeps it reliable.

Your weekly 20-minute review

1. **Open the Overview pane.** Clear anything waiting on you.
2. **Read three outputs at random,** not just the ones that look wrong.
3. **Fix the cause, not the output.** If a reply was off, update the skill or instructions and tell Claude to remember the correction. Fixing one draft fixes one draft; fixing the skill fixes all of them.
4. **Check your usage** in your account settings. If you're running out, move simple jobs to a smaller model or lower effort.
5. **Write down hours saved.** Compare against the baseline from Chapter 2.

Earn autonomy in stages

Treat every job like a new hire's probation:

- **Stage 1, Draft:** the agent drafts, you approve everything.
- **Stage 2, Act on low-risk work:** after two clean weeks, let it act on the safest category (sorting, internal reports, FAQ answers).
- **Stage 3, Report by exception:** it handles the routine work and only brings you exceptions.

Never skip to Stage 3 for anything involving money, customers in a bad moment, or public posts.

Guardrails that protect you

- **Write the approval list into your instructions** (Chapter 3). It's the single most important paragraph you'll write.
- **Never paste passwords, API keys or full card numbers** into chats, files or memory.
- **Keep sensitive data out of projects that don't need it.** Separate projects keep access separate.
- **Pause when something looks wrong.** **Project settings > General** has a **Pause** button that stops every thread and routine in the project at once.
- **Know your data rules.** If you handle health, legal, financial or children's data, check your obligations before you connect those systems.

The five most common mistakes

1. **Automating chaos.** If a process is a mess with humans, it'll be a faster mess with agents. Write the SOP first.
2. **Six departments at once.** Set up one, get it right, then copy what worked.
3. **No examples.** One real example of great work in a skill beats a page of rules.

4. **Skipping the review.** Most problems show up in week two, not day one.
5. **Giving every connector to every job.** Least access, always.

Your move

Put a recurring 20-minute “AI workforce review” on your calendar for the same time every week. Then, if you want, ask your project to remind you: make it a routine that sends you a list of what’s waiting on you every Friday.

Your 30-day rollout plan

Week 1: Foundation. Map your work (Chapter 2). Choose your first department and three jobs. Create the project, write the instructions, connect only the apps you need. Send one small task a day and correct what comes back.

Week 2: Procedures. Write a skill for each of your three jobs. Add one real example to each. Let the project propose threads, and approve them one at a time.

Week 3: Autopilot. Turn your most predictable job into a routine. Read every run. Move one job from “Draft” to “Act on low-risk work” if it’s earned it.

Week 4: Measure and expand. Compare hours spent against your baseline. Fix the weakest skill. Pick the next department and repeat from Week 1, reusing your voice skill and instructions.

By day 30 you should have one department running with less than an hour a week of your attention, and a pattern you can copy to every other part of the business.

Want help building yours?

This guide gives you the blueprint. If you'd rather have someone who has run operations for a 50-person company map your workflows, decide what to hand off first, and design the workforce with you, that's exactly what I do.

Book a free discovery call at virtualexpert.net. We'll look at where your week goes and which three jobs would give you the most time back.

And if your bottleneck is getting clients, start with **The AI Client Playbook** and **The AI Client Kit**: the step-by-step system and 110 ready-to-use prompts for finding, winning and keeping clients with AI.

Product names, menus and plan details reflect Anthropic's documentation as of October 2026 and may change. Claude is a product of Anthropic; this guide is independent and not affiliated with or endorsed by Anthropic.